

# Ensemble modelling

August 15, 2023

## 1 Dataset

<https://www.kaggle.com/datasets/kmader/skin-cancer-mnist-ham10000>

Cases include a representative collection of all important diagnostic categories in the realm of pigmented lesions: - Actinic keratoses and intraepithelial carcinoma / Bowen's disease (AKIEC), - basal cell carcinoma (BCC) - benign keratosis-like lesions (solar lentigines / seborrheic keratoses and lichen-planus like keratoses, BKL) - dermatofibroma (DF) - melanoma (MEL) - melanocytic nevi (NV) - vascular lesions (angiomas, angiokeratomas, pyogenic granulomas and hemorrhage, VASC).

The classes and the target labels have been listed below

Category	Class ID
melanoma (MEL)	0
melanocytic nevi (NV)	1
basal cell carcinoma (BCC)	2
Actinic keratoses and intraepithelial carcinoma / Bowen's disease (AKIEC)	3
benign keratosis-like lesions (solar lentigines / seborrheic keratoses and lichen-planus like keratoses, BKL)	4
dermatofibroma (DF)	5
vascular lesions (angiomas, angiokeratomas, pyogenic granulomas and hemorrhage, VASC)	6

## 2 Classification features

The LBP calculated from the images are being used as features for image classification.

## 3 Untransformed

Generating classification report and confusion matrix of images that were not log transformed before segmentation

### 3.1 SVM classification


Train test split is 80% and 20% respectively.

### 3.1.1 10-fold cross validation

Fold 1: Accuracy = 0.67  
Fold 2: Accuracy = 0.67  
Fold 3: Accuracy = 0.68  
Fold 4: Accuracy = 0.67  
Fold 5: Accuracy = 0.67  
Fold 6: Accuracy = 0.68  
Fold 7: Accuracy = 0.68  
Fold 8: Accuracy = 0.66  
Fold 9: Accuracy = 0.66  
Fold 10: Accuracy = 0.67  
Mean Accuracy: 0.67

[8]: SVC()

### 3.1.2 Confusion Matrix



### 3.1.3 Classification Report

Sensitivity: 1.00

Specificity: 0.00

	precision	recall	f1-score	support
0	0.00	0.00	0.00	239
1	0.67	1.00	0.80	1338
2	0.00	0.00	0.00	106
3	0.00	0.00	0.00	62
4	0.00	0.00	0.00	206
5	0.00	0.00	0.00	22
6	0.00	0.00	0.00	30
accuracy			0.67	2003
macro avg	0.10	0.14	0.11	2003
weighted avg	0.45	0.67	0.54	2003

```
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))  
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))  
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))
```

## 3.2 Random forrest classifier

Train test split is 80% and 20% respectively.


### 3.2.1 10-fold cross validation

Fold 1: Accuracy = 0.68  
Fold 2: Accuracy = 0.68  
Fold 3: Accuracy = 0.69  
Fold 4: Accuracy = 0.68  
Fold 5: Accuracy = 0.69  
Fold 6: Accuracy = 0.69  
Fold 7: Accuracy = 0.70  
Fold 8: Accuracy = 0.67

Fold 9: Accuracy = 0.65  
 Fold 10: Accuracy = 0.67  
 Mean Accuracy: 0.68

[11]: RandomForestClassifier()

### 3.2.2 Confusion matrix



### 3.2.3 Classification Report

Sensitivity: 0.99

Specificity: 0.09

	precision	recall	f1-score	support
0	0.40	0.08	0.13	239
1	0.73	0.96	0.83	1338
2	0.32	0.11	0.17	106
3	0.32	0.26	0.29	62
4	0.29	0.13	0.18	206

5	0.00	0.00	0.00	22
6	0.00	0.00	0.00	30
accuracy			0.68	2003
macro avg	0.29	0.22	0.23	2003
weighted avg	0.59	0.68	0.60	2003

```

/home/rajdeep/projects/en/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/home/rajdeep/projects/en/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/home/rajdeep/projects/en/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

### 3.3 Adaboost

Train test split is 80% and 20% respectively.

#### 3.3.1 10-fold cross validation


```

Fold 1: Accuracy = 0.60
Fold 2: Accuracy = 0.59
Fold 3: Accuracy = 0.68
Fold 4: Accuracy = 0.62
Fold 5: Accuracy = 0.68
Fold 6: Accuracy = 0.67
Fold 7: Accuracy = 0.55
Fold 8: Accuracy = 0.62
Fold 9: Accuracy = 0.64
Fold 10: Accuracy = 0.64
Mean Accuracy: 0.63

```

```
[14]: AdaBoostClassifier()
```

### 3.3.2 Confusion Matrix



### 3.3.3 Classification Report

Sensitivity: 1.00

Specificity: 0.00

	precision	recall	f1-score	support
0	1.00	0.00	0.01	239
1	0.71	0.96	0.82	1338
2	0.22	0.08	0.11	106
3	0.20	0.26	0.22	62
4	0.33	0.05	0.09	206
5	0.00	0.00	0.00	22
6	0.09	0.07	0.08	30
accuracy			0.66	2003
macro avg	0.36	0.20	0.19	2003
weighted avg	0.65	0.66	0.57	2003

## 4 Transformed

Generating classification report and confusion matrix of images that were log transformed before segmentation

### 4.1 SVM classification


Train test split is 80% and 20% respectively.

#### 4.1.1 10-fold cross validation

```
Fold 1: Accuracy = 0.67
Fold 2: Accuracy = 0.67
Fold 3: Accuracy = 0.68
Fold 4: Accuracy = 0.67
Fold 5: Accuracy = 0.67
Fold 6: Accuracy = 0.68
Fold 7: Accuracy = 0.68
Fold 8: Accuracy = 0.66
Fold 9: Accuracy = 0.66
Fold 10: Accuracy = 0.67
Mean Accuracy: 0.67
```

```
[23]: SVC()
```

#### 4.1.2 confusion Matrix



#### 4.1.3 Classification Report

Sensitivity: 1.00

Specificity: 0.00

	precision	recall	f1-score	support
0	0.00	0.00	0.00	239
1	0.67	1.00	0.80	1338
2	0.00	0.00	0.00	106
3	0.00	0.00	0.00	62
4	0.00	0.00	0.00	206
5	0.00	0.00	0.00	22
6	0.00	0.00	0.00	30
accuracy			0.67	2003
macro avg	0.10	0.14	0.11	2003
weighted avg	0.45	0.67	0.54	2003




```
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))
```

## 4.2 Random forrest classifier

### 4.2.1 10-fold cross validation

```
Fold 1: Accuracy = 0.68  
Fold 2: Accuracy = 0.69  
Fold 3: Accuracy = 0.70  
Fold 4: Accuracy = 0.68  
Fold 5: Accuracy = 0.69  
Fold 6: Accuracy = 0.68  
Fold 7: Accuracy = 0.69  
Fold 8: Accuracy = 0.68  
Fold 9: Accuracy = 0.68  
Fold 10: Accuracy = 0.67  
Mean Accuracy: 0.68
```

### 4.2.2 Confusion Matrix



### 4.2.3 Classification Report

Sensitivity: 0.99

Specificity: 0.12

	precision	recall	f1-score	support
0	0.45	0.10	0.17	239
1	0.73	0.96	0.83	1338
2	0.26	0.11	0.16	106
3	0.37	0.26	0.30	62
4	0.30	0.13	0.18	206
5	0.00	0.00	0.00	22
6	0.00	0.00	0.00	30
accuracy			0.68	2003
macro avg	0.30	0.22	0.23	2003
weighted avg	0.60	0.68	0.61	2003

```

/home/rajdeep/projects/en/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/home/rajdeep/projects/en/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
/home/rajdeep/projects/en/lib/python3.10/site-
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

### 4.3 Adaboost

Train test split is 80% and 20% respectively.

#### 4.3.1 10-fold cross validation


```

Fold 1: Accuracy = 0.65
Fold 2: Accuracy = 0.66
Fold 3: Accuracy = 0.66
Fold 4: Accuracy = 0.67
Fold 5: Accuracy = 0.65
Fold 6: Accuracy = 0.66
Fold 7: Accuracy = 0.64
Fold 8: Accuracy = 0.64
Fold 9: Accuracy = 0.65
Fold 10: Accuracy = 0.63
Mean Accuracy: 0.65

```

```
[29]: AdaBoostClassifier()
```

### 4.3.2 Confusion Matrix



### 4.3.3 Classification report

Sensitivity: 0.99

Specificity: 0.04

	precision	recall	f1-score	support
0	0.40	0.03	0.06	239
1	0.69	0.95	0.80	1338
2	0.11	0.03	0.05	106
3	0.22	0.18	0.19	62
4	0.29	0.03	0.05	206
5	0.00	0.00	0.00	22
6	0.04	0.07	0.05	30
accuracy			0.65	2003
macro avg	0.25	0.18	0.17	2003
weighted avg	0.55	0.65	0.56	2003

```
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))  
/home/rajdeep/projects/en/lib/python3.10/site-  
packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning:  
Precision and F-score are ill-defined and being set to 0.0 in labels with no  
predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, msg_start, len(result))
```